

Titanic - Machine Learning from Disaster

December 21, 2023

Initially importing usefull liraries

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

Importing training data set in a variable

```
[2]: titanic_train = pd.read_csv(r'C:\Users\Bikash shah\Desktop\titanic\train.csv')
```

```
[3]: titanic_train
```

```
[3]:      PassengerId  Survived  Pclass  \
0                1         0        3
1                2         1        1
2                3         1        3
3                4         1        1
4                5         0        3
..          ...     ...     ...
886           887         0        2
887           888         1        1
888           889         0        3
889           890         1        1
890           891         0        3
```

```
      Name      Sex  Age  SibSp  \
0  Braund, Mr. Owen Harris    male  22.0    1
1  Cumings, Mrs. John Bradley (Florence Briggs Th...  female  38.0    1
2  Heikkinen, Miss. Laina    female  26.0    0
3  Futrelle, Mrs. Jacques Heath (Lily May Peel)    female  35.0    1
4  Allen, Mr. William Henry    male  35.0    0
..          ...     ...     ...
886  Montvila, Rev. Juozas    male  27.0    0
887  Graham, Miss. Margaret Edith    female  19.0    0
888  Johnston, Miss. Catherine Helen "Carrie"    female   NaN    1
889  Behr, Mr. Karl Howell    male  26.0    0
890  Dooley, Mr. Patrick    male  32.0    0
```

```
      Parch      Ticket    Fare Cabin Embarked
```

0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	NaN	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	NaN	S
..
886	0	211536	13.0000	NaN	S
887	0	112053	30.0000	B42	S
888	2	W./C. 6607	23.4500	NaN	S
889	0	111369	30.0000	C148	C
890	0	370376	7.7500	NaN	Q

[891 rows x 12 columns]

```
[4]: titanic_train.head()
```

```
[4]: PassengerId  Survived  Pclass  \
0            1         0         3
1            2         1         1
2            3         1         3
3            4         1         1
4            5         0         3
```

	Name	Sex	Age	SibSp	\
0	Braund, Mr. Owen Harris	male	22.0	1	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	
2	Heikkinen, Miss. Laina	female	26.0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	
4	Allen, Mr. William Henry	male	35.0	0	

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	NaN	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	NaN	S


Looking for a correlation among the variables at the datasets

```
[5]: import seaborn as sns

sns.heatmap(titanic_train.corr(), cmap='YlGnBu')
plt.show()
```

C:\Users\Bikash shah\AppData\Local\Temp\ipykernel_8012\1610899979.py:3:
FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
sns.heatmap(titanic_train.corr(), cmap='YlGnBu')
```



Splitting this training data sets into two data set which is training and testing with the help of StratifiedShuffleSplit in the same ratio of survived, pclass and sex variables.

```
[6]: from sklearn.model_selection import StratifiedShuffleSplit

split = StratifiedShuffleSplit(n_splits=1, test_size=0.2)
for train_indicates, test_indicates in split.split(titanic_train,
    ↪titanic_train[['Survived', 'Pclass', 'Sex']]):
    strat_train_set = titanic_train.loc[train_indicates]
    strat_test_set = titanic_train.loc[test_indicates]
```


Looking in histogram is the data set distributed rationally in both training and testing data sets.

```
[7]: plt.subplot(1,2,1)
strat_train_set['Survived'].hist()
strat_train_set['Pclass'].hist()

plt.subplot(1,2,2)
```

```
strat_test_set['Survived'].hist()
strat_test_set['Pclass'].hist()

plt.show()
```



We will look for the information about having null value or not.

```
[8]: strat_train_set.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 712 entries, 444 to 171
Data columns (total 12 columns):
#   Column      Non-Null Count  Dtype
---  -
0   PassengerId  712 non-null    int64
1   Survived     712 non-null    int64
2   Pclass       712 non-null    int64
3   Name         712 non-null    object
4   Sex          712 non-null    object
5   Age         570 non-null    float64
6   SibSp        712 non-null    int64
7   Parch        712 non-null    int64
8   Ticket       712 non-null    object
```

```
9   Fare           712 non-null    float64
10  Cabin          161 non-null    object
11  Embarked       710 non-null    object
dtypes: float64(2), int64(5), object(5)
memory usage: 72.3+ KB
```

1 Creating a pipeline for cleaning the data set.

We will create a class for adding mean value at Age variable for the null value.

```
[9]: from sklearn.base import BaseEstimator, TransformerMixin
     from sklearn.impute import SimpleImputer

     class AgeImputer(BaseEstimator, TransformerMixin):

         def fit(self, X, y=None):
             return self

         def transform(self, X):
             imputer = SimpleImputer(strategy = 'mean')
             X['Age'] = imputer.fit_transform(X[['Age']])
             return X
```

We will create a class for changing the categorical data into binary data which 0 and 1.

```
[10]: from sklearn.preprocessing import OneHotEncoder

     class FeatureEncoder(BaseEstimator, TransformerMixin):

         def fit(self, X, y=None):
             return self

         def transform(self, X):
             encoder = OneHotEncoder()
             matrix = encoder.fit_transform(X[['Embarked']]).toarray()

             columne_names = ['C', 'S', 'Q', 'N']

             for i in range(len(matrix.T)):
                 X[columne_names[i]] = matrix.T[i]

             matrix = encoder.fit_transform(X[['Sex']]).toarray()

             columne_names = ['Female', 'Male']

             for i in range(len(matrix.T)):
                 X[columne_names[i]] = matrix.T[i]
```

```
return X
```

We will create a class for removing the variable which we are not going to need.

```
[11]: class FeatureDropper(BaseEstimator, TransformerMixin):

    def fit(self, X, y=None):
        return self

    def transform(self, X):
        return X.drop(['Embarked', 'Name', 'Ticket', 'Cabin', 'Sex', 'N'],
            ↪axis=1, errors = 'ignore')
```

We have used pipeline to create a pipeline to make data set into numeric data set which will help us predicting more accurately.

```
[12]: from sklearn.pipeline import Pipeline

pipeline = Pipeline([('ageimputer', AgeImputer()),
                      ('featureencoder', FeatureEncoder()),
                      ('featuredropper', FeatureDropper())])
```

Changing the training data set into numeric data set.

```
[13]: strat_train_set = pipeline.fit_transform(strat_train_set)
```

```
[14]: strat_train_set.head()
```

```
[14]:
```

	PassengerId	Survived	Pclass	Age	SibSp	Parch	Fare	C	S	\
444	445	1	3	29.581	0	0	8.1125	0.0	0.0	
592	593	0	3	47.000	0	0	7.2500	0.0	0.0	
580	581	1	2	25.000	1	1	30.0000	0.0	0.0	
46	47	0	3	29.581	1	0	15.5000	0.0	1.0	
852	853	0	3	9.000	1	1	15.2458	1.0	0.0	

	Q	Female	Male
444	1.0	0.0	1.0
592	1.0	0.0	1.0
580	1.0	1.0	0.0
46	0.0	0.0	1.0
852	0.0	1.0	0.0

```
[15]: strat_train_set.describe()
```

```
[15]:
```

	PassengerId	Survived	Pclass	Age	SibSp	\
count	712.000000	712.000000	712.000000	712.000000	712.000000	
mean	445.676966	0.384831	2.308989	29.581000	0.518258	
std	259.409965	0.486897	0.835249	13.088319	1.086785	
min	1.000000	0.000000	1.000000	0.420000	0.000000	

25%	220.500000	0.000000	2.000000	22.000000	0.000000
50%	443.500000	0.000000	3.000000	29.581000	0.000000
75%	675.250000	1.000000	3.000000	35.000000	1.000000
max	890.000000	1.000000	3.000000	80.000000	8.000000

	Parch	Fare	C	S	Q	Female \
count	712.000000	712.000000	712.000000	712.000000	712.000000	712.000000
mean	0.356742	31.632689	0.181180	0.089888	0.726124	0.351124
std	0.764349	47.914117	0.385438	0.286222	0.446260	0.477657
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	7.895800	0.000000	0.000000	0.000000	0.000000
50%	0.000000	14.454200	0.000000	0.000000	1.000000	0.000000
75%	0.000000	30.178100	0.000000	0.000000	1.000000	1.000000
max	5.000000	512.329200	1.000000	1.000000	1.000000	1.000000

	Male
count	712.000000
mean	0.648876
std	0.477657
min	0.000000
25%	0.000000
50%	1.000000
75%	1.000000
max	1.000000

```
[16]: strat_train_set.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 712 entries, 444 to 171
Data columns (total 12 columns):
#   Column      Non-Null Count  Dtype
---  -
0   PassengerId  712 non-null    int64
1   Survived     712 non-null    int64
2   Pclass       712 non-null    int64
3   Age          712 non-null    float64
4   SibSp        712 non-null    int64
5   Parch        712 non-null    int64
6   Fare         712 non-null    float64
7   C            712 non-null    float64
8   S            712 non-null    float64
9   Q            712 non-null    float64
10  Female       712 non-null    float64
11  Male         712 non-null    float64
dtypes: float64(7), int64(5)
memory usage: 72.3 KB
```

Changing data set into data points for training data set.

```
[17]: from sklearn.preprocessing import StandardScaler

X = strat_train_set.drop(['Survived'], axis=1)
y = strat_train_set['Survived']

scaler = StandardScaler()
X_data = scaler.fit_transform(X)
y_data = y.to_numpy()
```

Applying RandomForestClassifier model algorithm for predicting the testing set.

```
[18]: from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV

clf = RandomForestClassifier()

param_grid = [{'n_estimators': [10, 100, 200, 500], 'max_depth': [None, 5, 10],
               ↪ 'min_samples_split': [2, 3, 4]}]

grid_search = GridSearchCV(clf, param_grid, cv=3, scoring='accuracy',
                           ↪ return_train_score=True)
grid_search.fit(X_data, y_data)
```

```
[18]: GridSearchCV(cv=3, estimator=RandomForestClassifier(),
                  param_grid=[{'max_depth': [None, 5, 10],
                               'min_samples_split': [2, 3, 4],
                               'n_estimators': [10, 100, 200, 500]}],
                  return_train_score=True, scoring='accuracy')
```

```
[19]: final_clf = grid_search.best_estimator_
```

```
[20]: final_clf
```

```
[20]: RandomForestClassifier(min_samples_split=4, n_estimators=500)
```

As we can see that max depth the tree went up to is 10, minimum split is 3, and estimator is 500. Now we are going to do same thing for testing data set.

```
[21]: strat_test_set = pipeline.fit_transform(strat_test_set)
```

```
[22]: X_test = strat_test_set.drop(['Survived'], axis=1)
y_test = strat_test_set['Survived']

scaler = StandardScaler()
X_data_test = scaler.fit_transform(X_test)
y_data_test = y_test.to_numpy()
```

```
[23]: final_clf.score(X_data_test, y_data_test)
```



```
[23]: 0.7988826815642458
```

```
[24]: final_data = pipeline.fit_transform(titanic_train)
```

```
[25]: final_data.head()
```

```
[25]:
```

	PassengerId	Survived	Pclass	Age	SibSp	Parch	Fare	C	S	Q	\
0	1	0	3	22.0	1	0	7.2500	0.0	0.0	1.0	
1	2	1	1	38.0	1	0	71.2833	1.0	0.0	0.0	
2	3	1	3	26.0	0	0	7.9250	0.0	0.0	1.0	
3	4	1	1	35.0	1	0	53.1000	0.0	0.0	1.0	
4	5	0	3	35.0	0	0	8.0500	0.0	0.0	1.0	

	Female	Male
0	0.0	1.0
1	1.0	0.0
2	1.0	0.0
3	1.0	0.0
4	0.0	1.0

Now we are going to use the real whole training data for in the algorithm.

```
[26]: X_final = final_data.drop(['Survived'], axis=1)
      y_final = final_data['Survived']

      scaler = StandardScaler()
      X_data_final = scaler.fit_transform(X_final)
      y_data_final = y_final.to_numpy()
```

```
[27]: prod_clf = RandomForestClassifier()

      param_grid = [{'n_estimators': [10, 100, 200, 500], 'max_depth': [None, 5, 10],
      ↪ 'min_samples_split': [2, 3, 4]}]

      grid_search = GridSearchCV(prod_clf, param_grid, cv=3, scoring='accuracy',
      ↪ return_train_score=True)
      grid_search.fit(X_data_final, y_data_final)
```

```
[27]: GridSearchCV(cv=3, estimator=RandomForestClassifier(),
      param_grid=[{'max_depth': [None, 5, 10],
      'min_samples_split': [2, 3, 4],
      'n_estimators': [10, 100, 200, 500]}],
      return_train_score=True, scoring='accuracy')
```

```
[28]: prod_final_clf = grid_search.best_estimator_
```

```
[29]: prod_final_clf
```

```
[29]: RandomForestClassifier(max_depth=5, min_samples_split=3, n_estimators=200)
```

With the help of model we have trained for training data set, we are going to predict the testing data by repeating the same process.

```
[30]: titanic_test = pd.read_csv(r'C:\Users\Bikash shah\Desktop\titanic\test.csv')
```

```
[31]: titanic_test
```

```
[31]:
```

	PassengerId	Pclass	Name \
0	892	3	Kelly, Mr. James
1	893	3	Wilkes, Mrs. James (Ellen Needs)
2	894	2	Myles, Mr. Thomas Francis
3	895	3	Wirz, Mr. Albert
4	896	3	Hirvonen, Mrs. Alexander (Helga E Lindqvist)
..
413	1305	3	Spector, Mr. Woolf
414	1306	1	Oliva y Ocana, Dona. Fermina
415	1307	3	Saether, Mr. Simon Sivertsen
416	1308	3	Ware, Mr. Frederick
417	1309	3	Peter, Master. Michael J

	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	male	34.5	0	0	330911	7.8292	NaN	Q
1	female	47.0	1	0	363272	7.0000	NaN	S
2	male	62.0	0	0	240276	9.6875	NaN	Q
3	male	27.0	0	0	315154	8.6625	NaN	S
4	female	22.0	1	1	3101298	12.2875	NaN	S
..
413	male	NaN	0	0	A.5. 3236	8.0500	NaN	S
414	female	39.0	0	0	PC 17758	108.9000	C105	C
415	male	38.5	0	0	SOTON/O.Q. 3101262	7.2500	NaN	S
416	male	NaN	0	0	359309	8.0500	NaN	S
417	male	NaN	1	1	2668	22.3583	NaN	C

[418 rows x 11 columns]

```
[32]: final_test_data = pipeline.fit_transform(titanic_test)
```

```
[33]: final_test_data
```

```
[33]:
```

	PassengerId	Pclass	Age	SibSp	Parch	Fare	C	S	Q	\
0	892	3	34.50000	0	0	7.8292	0.0	1.0	0.0	
1	893	3	47.00000	1	0	7.0000	0.0	0.0	1.0	
2	894	2	62.00000	0	0	9.6875	0.0	1.0	0.0	
3	895	3	27.00000	0	0	8.6625	0.0	0.0	1.0	
4	896	3	22.00000	1	1	12.2875	0.0	0.0	1.0	
..	

413	1305	3	30.27259	0	0	8.0500	0.0	0.0	1.0
414	1306	1	39.00000	0	0	108.9000	1.0	0.0	0.0
415	1307	3	38.50000	0	0	7.2500	0.0	0.0	1.0
416	1308	3	30.27259	0	0	8.0500	0.0	0.0	1.0
417	1309	3	30.27259	1	1	22.3583	1.0	0.0	0.0

	Female	Male
0	0.0	1.0
1	1.0	0.0
2	0.0	1.0
3	0.0	1.0
4	1.0	0.0
..
413	0.0	1.0
414	1.0	0.0
415	0.0	1.0
416	0.0	1.0
417	0.0	1.0

[418 rows x 11 columns]

```
[34]: X_final_test = final_test_data
X_final_test = X_final_test.fillna(method='ffill')

scaler = StandardScaler()
X_data_final_test = scaler.fit_transform(X_final_test)
```

```
[35]: X_data_final_test
```

```
[35]: array([[ -1.72791209,  0.87348191,  0.3349926 , ..., -1.35067551,
          -0.75592895,  0.75592895],
        [ -1.71962474,  0.87348191,  1.32553003, ...,  0.74037028,
          1.32287566, -1.32287566],
        [ -1.71133739, -0.31581919,  2.51417495, ..., -1.35067551,
          -0.75592895,  0.75592895],
        ...,
        [  1.71133739,  0.87348191,  0.65196458, ...,  0.74037028,
          -0.75592895,  0.75592895],
        [  1.71962474,  0.87348191,  0.          , ...,  0.74037028,
          -0.75592895,  0.75592895],
        [  1.72791209,  0.87348191,  0.          , ..., -1.35067551,
          -0.75592895,  0.75592895]])
```

Now we are going to predict the final test variable with the help of final classifier.

```
[36]: predictions = prod_final_clf.predict(X_data_final_test)
```

```
[37]: predictions
```

```
[37]: array([0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 1, 0, 1, 1, 0, 0, 1, 1, 0, 0,
            1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1,
            1, 0, 0, 0, 1, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1,
            1, 0, 0, 1, 1, 0, 1, 0, 1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 1, 1,
            1, 1, 1, 0, 1, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0,
            0, 1, 1, 1, 1, 0, 0, 1, 0, 1, 1, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0,
            0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1,
            0, 0, 1, 1, 0, 1, 1, 1, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1,
            1, 0, 1, 1, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 1, 1,
            0, 1, 1, 1, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0,
            1, 0, 1, 0, 1, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1,
            0, 0, 0, 0, 1, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1,
            0, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 0, 0, 1, 1, 1, 1, 0,
            0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1,
            0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0,
            1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 1, 1, 0,
            0, 0, 1, 0, 1, 0, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0,
            1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1,
            0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 0, 1, 0, 0, 0],
            dtype=int64)
```

```
[38]: print(len(predictions))
```

418

We will write this predicted data into CSV file.

```
[100]: final_df = pd.DataFrame(titanic_test['PassengerId'])
        final_df['Survived'] = predictions
        final_df.to_csv(r'C:\Users\Bikash shah\Desktop\titanic\prediction.csv',
                        index=False)
```

```
[ ]:
```